


Mixed-Precision Profile Map — arctic-embed-m
(orange = INT8, grey = FP16, blue = FP32)

★ Pareto optimal



† Token Embedding (`nn.Embedding`) is not an `nn.Linear` layer — `quantize_dynamic` leaves it in FP32 for all INT8 profiles.